

FIN 580 Homework 2 - Tejas Dhomne | UIN - 661586178

Import pandas_datareader as pdr

```
In [1]: import pandas as pd
import pandas_datareader as pdr
import numpy as np
```

Question 1

1.1) Create a new column in df1 named month1 that extracts month information from the index.

```
In [2]: #1.1
df1 = pdr.get_data_yahoo('GOOGL', start='2020-01-01', end='2020-12-31')
df1
```

Out[2]:

| | High | Low | Open | Close | Volume | Adj Close |
|------------|-----------|-----------|-----------|-----------|----------|-----------|
| Date | | | | | | |
| 2020-01-02 | 68.433998 | 67.324501 | 67.420502 | 68.433998 | 27278000 | 68.433998 |
| 2020-01-03 | 68.687500 | 67.365997 | 67.400002 | 68.075996 | 23408000 | 68.075996 |
| 2020-01-06 | 69.916000 | 67.550003 | 67.581497 | 69.890503 | 46768000 | 69.890503 |
| 2020-01-07 | 70.175003 | 69.578003 | 70.023003 | 69.755501 | 34330000 | 69.755501 |
| 2020-01-08 | 70.592499 | 69.631500 | 69.740997 | 70.251999 | 35314000 | 70.251999 |
| ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 87.120499 | 86.217499 | 86.449997 | 86.708000 | 9312000 | 86.708000 |
| 2020-12-28 | 89.349998 | 87.091003 | 87.245499 | 88.697998 | 27650000 | 88.697998 |
| 2020-12-29 | 89.423500 | 87.755501 | 89.361504 | 87.888000 | 19726000 | 87.888000 |
| 2020-12-30 | 88.388000 | 86.400002 | 88.250000 | 86.812500 | 21026000 | 86.812500 |
| 2020-12-31 | 87.875000 | 86.804497 | 86.863503 | 87.632004 | 21070000 | 87.632004 |

253 rows × 6 columns

```
In [3]: df1['month1']=df1.index.month
df1
```

Out[3]:

| | High | Low | Open | Close | Volume | Adj Close | month1 |
|------------|-----------|-----------|-----------|-----------|----------|-----------|--------|
| Date | | | | | | | |
| 2020-01-02 | 68.433998 | 67.324501 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 |
| 2020-01-03 | 68.687500 | 67.365997 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 |
| 2020-01-06 | 69.916000 | 67.550003 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 |
| 2020-01-07 | 70.175003 | 69.578003 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 |
| 2020-01-08 | 70.592499 | 69.631500 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 87.120499 | 86.217499 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 |
| 2020-12-28 | 89.349998 | 87.091003 | 87.245499 | 88.697998 | 27650000 | 88.697998 | 12 |
| 2020-12-29 | 89.423500 | 87.755501 | 89.361504 | 87.888000 | 19726000 | 87.888000 | 12 |
| 2020-12-30 | 88.388000 | 86.400002 | 88.250000 | 86.812500 | 21026000 | 86.812500 | 12 |

1.2) Use replace() to create a new column in df1 named month2 that replaces integers in month1 with abbreviations of the names of the months. For example, replace 1 with Jan, 2 with Feb, 3 with Mar, 4 with Apr, 5 with May, 6 with Jun, 7 with Jul, 8 with Aug, 9 with Sep, 10 with Oct, 11 with Nov, and 12 with Dec. Do not modify the numeric values in month1.

In [4]:

```
ch1"].replace({1:"Jan",2:"Feb",3:"Mar",4:"Apr",5:"May",6:"Jun",7:"Jul",8:"Aug",9:"Sep",10:"Oct",11:"Nov",12:"Dec"})
```

Out[4]:

| | High | Low | Open | Close | Volume | Adj Close | month1 | month2 |
|------------|-----------|-----------|-----------|-----------|----------|-----------|--------|--------|
| Date | | | | | | | | |
| 2020-01-02 | 68.433998 | 67.324501 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 | Jan |
| 2020-01-03 | 68.687500 | 67.365997 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 | Jan |
| 2020-01-06 | 69.916000 | 67.550003 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 | Jan |
| 2020-01-07 | 70.175003 | 69.578003 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 | Jan |
| 2020-01-08 | 70.592499 | 69.631500 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 | Jan |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 87.120499 | 86.217499 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 | Dec |
| 2020-12-28 | 89.349998 | 87.091003 | 87.245499 | 88.697998 | 27650000 | 88.697998 | 12 | Dec |
| 2020-12-29 | 89.423500 | 87.755501 | 89.361504 | 87.888000 | 19726000 | 87.888000 | 12 | Dec |
| 2020-12-30 | 88.388000 | 86.400002 | 88.250000 | 86.812500 | 21026000 | 86.812500 | 12 | Dec |
| 2020-12-31 | 87.875000 | 86.804497 | 86.863503 | 87.632004 | 21070000 | 87.632004 | 12 | Dec |

253 rows × 8 columns

1.3) Return the counts of unique values in month2 and sort the series of value counts by the counts in an ascending order.

In [5]:

```
#1.3
df1["month2"].value_counts().sort_values()
```

Out[5]:

```
Feb    19
May    20
Nov    20
Jan    21
Apr    21
Aug    21
Sep    21
Mar    22
Jun    22
Jul    22
Oct    22
Dec    22
Name: month2, dtype: int64
```

1.4) Create a new column in df1 named high_minus_low1 that represents the difference between High and Low.

```
In [6]: df1['high_minus_low1']=df1["High"]-df1["Low"]
df1
```

| Date | | | | | | | | |
|------------|-----------|-----------|-----------|-----------|----------|-----------|-----|-----|
| 2020-01-02 | 68.433998 | 67.324501 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 | Jan |
| 2020-01-03 | 68.687500 | 67.365997 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 | Jan |
| 2020-01-06 | 69.916000 | 67.550003 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 | Jan |
| 2020-01-07 | 70.175003 | 69.578003 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 | Jan |
| 2020-01-08 | 70.592499 | 69.631500 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 | Jan |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 87.120499 | 86.217499 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 | Dec |
| 2020-12-28 | 89.349998 | 87.091003 | 87.245499 | 88.697998 | 27650000 | 88.697998 | 12 | Dec |

1.5) Create a function named dollars with one parameter. The function rounds the argument to two decimal places, converts the number to a string, and adds the string to “ dollars”. After creating the function, dollars(22.190063) should return ‘22.19 dollars’.

```
In [7]: #1.5
def dollars(x):
    return str(round(x,2))+ " dollars"
```

```
In [8]: dollars(22.190063)
```

```
Out[8]: '22.19 dollars'
```

1.6) Apply the dollars function to all the values in the high_minus_low1 column. Create a new column in df1 named high_minus_low2 to store the result. Do not modify the values in high_minus_low1.

```
In [9]: #1.6
df1['high_minus_low2']=df1["high_minus_low1"].apply(dollars)
df1
```

| Date | | | | | | | | | |
|------------|-----------|-----------|-----------|-----------|----------|-----------|-----|-----|-----|
| 2020-01-02 | 68.433998 | 67.324501 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 | Jan | |
| 2020-01-03 | 68.687500 | 67.365997 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 | Jan | |
| 2020-01-06 | 69.916000 | 67.550003 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 | Jan | |
| 2020-01-07 | 70.175003 | 69.578003 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 | Jan | |
| 2020-01-08 | 70.592499 | 69.631500 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 | Jan | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 87.120499 | 86.217499 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 | Dec | |
| 2020-12-28 | 89.349998 | 87.091003 | 87.245499 | 88.697998 | 27650000 | 88.697998 | 12 | Dec | |

1.7) Use a lambda function to combine steps used in questions 1.5 and 1.6 into one step. Create a column in df1 named high_minus_low3 to store the result. Do not modify the values in high_minus_low1.

```
In [10]: #1.7
df1['high_minus_low3']=df1["high_minus_low1"].apply(lambda x: str(round(x,2)) + "
df1
```

Out[10]:

| | High | Low | Open | Close | Volume | Adj Close | month1 | month2 | high_ |
|------------|-----------|-----------|-----------|-----------|----------|-----------|--------|--------|-------|
| Date | | | | | | | | | |
| 2020-01-02 | 68.433998 | 67.324501 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 | Jan | |
| 2020-01-03 | 68.687500 | 67.365997 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 | Jan | |
| 2020-01-06 | 69.916000 | 67.550003 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 | Jan | |
| 2020-01-07 | 70.175003 | 69.578003 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 | Jan | |
| 2020-01-08 | 70.592499 | 69.631500 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 | Jan | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2020- | 87.120499 | 86.217499 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 | Dec | |

1.8) Return a row from df1 if high_minus_low2 is not equal to high_minus_low3. Since the values in high_minus_low2 and high_minus_low3 are the same, it should return an empty dataframe.

```
In [11]: #1.8
df1[df1["high_minus_low2"]!=df1["high_minus_low3"]]
```

Out[11]:

| | High | Low | Open | Close | Volume | Adj Close | month1 | month2 | high_minus_low1 | high_minus_low2 | high_minus_low3 |
|------|------|-----|------|-------|--------|--------------|--------|--------|-----------------|-----------------|-----------------|
| Date | | | | | | | | | | | |
| | | | | | | | | | | | |

1.9) Use drop(labels=, axis=, inplace=) to drop the "High", "Low", "high_minus_low1", "high_minus_low2", "high_minus_low3" columns from df1 and modify df1 in place. Specify the column names in a list using the labels parameter. Use axis="columns" to drop columns and inplace=True to modify df1 in place.

```
In [12]: #1.9
df1.drop(labels=["High", "Low", "high_minus_low1", "high_minus_low2", "high_minus_low3"], axis="columns", inplace=True)
df1
```

Out[12]:

| | Open | Close | Volume | Adj Close | month1 | month2 |
|------------|-----------|-----------|----------|-----------|--------|--------|
| Date | | | | | | |
| 2020-01-02 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 | Jan |
| 2020-01-03 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 | Jan |
| 2020-01-06 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 | Jan |
| 2020-01-07 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 | Jan |
| 2020-01-08 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 | Jan |
| ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 | Dec |
| 2020-12-28 | 87.245499 | 88.697998 | 27650000 | 88.697998 | 12 | Dec |
| 2020-12-29 | 89.361504 | 87.888000 | 19726000 | 87.888000 | 12 | Dec |
| 2020-12-30 | 88.250000 | 86.812500 | 21026000 | 86.812500 | 12 | Dec |

1.10) Use cut() to create a new column in df1 named price_bins to convert continuous values of the Adj Close column to four bins with an equal number of observations. Use the labels parameter to specify the following labels in order: low, moderate, high, and very high. Use include_lowest=True to specify that the first interval should be left-inclusive.

```
In [13]: #1.10
l1=[0, 0.25, 0.5, 0.75,1]
l2=[]
```

```
In [14]: #1.10
for i in l1:
    l2.append(df1["Adj Close"].quantile(i))
l2
```

```
Out[14]: [52.70650100708008,
69.8584976196289,
73.95549774169922,
78.22949981689453,
91.24849700927734]
```

```
In [15]: #1.10
df1['price_bins']=pd.cut(df1["Adj Close"],l2,labels=["low","moderate","high","very high"],name="price_bins")
df1
```

| Date | | | | | | | | |
|------------|-----------|-----------|----------|-----------|-----|-----|-----------|-----|
| 2020-01-02 | 67.420502 | 68.433998 | 27278000 | 68.433998 | 1 | Jan | low | |
| 2020-01-03 | 67.400002 | 68.075996 | 23408000 | 68.075996 | 1 | Jan | low | |
| 2020-01-06 | 67.581497 | 69.890503 | 46768000 | 69.890503 | 1 | Jan | moderate | |
| 2020-01-07 | 70.023003 | 69.755501 | 34330000 | 69.755501 | 1 | Jan | low | |
| 2020-01-08 | 69.740997 | 70.251999 | 35314000 | 70.251999 | 1 | Jan | moderate | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2020-12-24 | 86.449997 | 86.708000 | 9312000 | 86.708000 | 12 | Dec | very high | |
| 2020-12-28 | 87.245499 | 88.697998 | 27650000 | 88.697998 | 12 | Dec | very high | |
| 2020-12-29 | 89.361504 | 87.888000 | 19726000 | 87.888000 | 12 | Dec | very high | |
| 2020-12-30 | 88.250000 | 86.812500 | 21026000 | 86.812500 | 12 | Dec | very high | |
| 2020-12-31 | 86.863503 | 87.632004 | 21070000 | 87.632004 | 12 | Dec | very high | |

253 rows × 7 columns

1.11) Return the counts of unique values in price_bins and sort the series of value counts by the index.

```
In [16]: #1.11
df1["price_bins"].value_counts().sort_index()
```

```
Out[16]: low          64
moderate    63
high        63
very high   63
Name: price_bins, dtype: int64
```

1.12) Return a row from df1 if the value in the price_bins column is null. It should return an empty dataframe.

```
In [17]: #1.12
df1[df1["price_bins"].isnull()]
```

Out[17]:

| | Open | Close | Volume | Adj Close | month1 | month2 | price_bins |
|------|------|-------|--------|-----------|--------|--------|------------|
| Date | | | | | | | |

1.13) Use crosstab() to return two-way frequency counts of month1 and price_bins. Use margins=True to add row and column subtotals.

```
In [18]: #1.13
pd.crosstab(df1["month1"],df1["price_bins"],margins=True)
```

Out[18]:

| price_bins | low | moderate | high | very high | All |
|------------|-----|----------|------|-----------|-----|
| month1 | | | | | |
| 1 | 3 | 14 | 4 | 0 | 21 |
| 2 | 4 | 5 | 10 | 0 | 19 |
| 3 | 22 | 0 | 0 | 0 | 22 |
| 4 | 21 | 0 | 0 | 0 | 21 |
| 5 | 12 | 8 | 0 | 0 | 20 |
| 6 | 2 | 20 | 0 | 0 | 22 |
| 7 | 0 | 2 | 19 | 1 | 22 |
| 8 | 0 | 2 | 11 | 8 | 21 |
| 9 | 0 | 9 | 8 | 4 | 21 |
| 10 | 0 | 3 | 11 | 8 | 22 |
| 11 | 0 | 0 | 0 | 20 | 20 |
| 12 | 0 | 0 | 0 | 22 | 22 |
| All | 64 | 63 | 63 | 63 | 253 |

#1.13 - December month in 2020 has the greatest number of adjusted closing prices that fall in the very high category.

Question 2

2.1) Set the random seed value to 0 and use np.random.randint() to generate 100 numbers from a uniform distribution over [-50, 50). Convert the data type to float and save the result in a NumPy array named a1.


```
In [19]: #2.1
np.random.seed(0)
a1=np.random.randint(-50,50,100).astype('float')
a1
```

```
Out[19]: array([ -6.,  -3.,  14.,  17.,  17., -41.,  33., -29., -14.,  37.,  20.,
        38.,  38., -38.,   8.,  15., -11.,  37.,  -4.,  38.,  31., -13.,
       -25.,  27.,  22., -41., -30.,  30.,  19.,  29.,  -3.,  14.,  32.,
        49.,  38.,  -1., -21., -31., -31., -36., -11., -18.,  15., -41.,
         7., -18., -19.,  24., -27., -15.,  25.,   5., -22., -16., -50.,
       -50., -14.,   3., -45., -12., -33.,  29., -46.,  -8.,   8., -19.,
       -49.,  15.,  -9.,   7., -15., -39.,  -4.,  32.,  41., -50., -36.,
        49.,   3., -38.,  -8.,  34.,  25.,  18., -44.,  18.,  -3., -47.,
        26.,   2.,  28., -35., -30.,  49.,   8., -27.,  29., -37.,  35.,
       -2.]
```

2.2) Set the random seed value to 0 and use np.random.choice() to create a NumPy array named a2 and populate it with 20 integers randomly selected from [0, 100) without replacement. Use the replace=False parameter to sample without replacement.

```
In [20]: #2.2
np.random.seed(0)
a2=np.random.choice(np.arange(0,100), size=20, replace=False)
a2
```

```
Out[20]: array([26, 86,  2, 55, 75, 93, 16, 73, 54, 95, 53, 92, 78, 13,  7, 30, 22,
        24, 33,  8])
```

2.3) Replace numbers in a1 at indices a2 with missing values.

```
In [21]: #2.3
a1[a2]=np.nan
a1
```

```
Out[21]: array([ -6.,  -3.,  nan,  17.,  17., -41.,  33.,  nan,  nan,  37.,  20.,
        38.,  38.,  nan,   8.,  15.,  nan,  37.,  -4.,  38.,  31., -13.,
        nan,  27.,  nan, -41.,  nan,  30.,  19.,  29.,  nan,  14.,  32.,
        nan,  38.,  -1., -21., -31., -31., -36., -11., -18.,  15., -41.,
         7., -18., -19.,  24., -27., -15.,  25.,   5., -22.,  nan,  nan,
        nan, -14.,   3., -45., -12., -33.,  29., -46.,  -8.,   8., -19.,
       -49.,  15.,  -9.,   7., -15., -39.,  -4.,  nan,  41.,  nan, -36.,
        49.,  nan, -38.,  -8.,  34.,  25.,  18., -44.,  18.,  nan, -47.,
        26.,   2.,  28., -35.,  nan,  nan,   8.,  nan,  29., -37.,  35.,
       -2.]
```

2.4) Use reshape() to convert a1 from a 100-element vector to a 25*4 matrix.

In [22]: #2.4

a1=a1.reshape(25,4)

a1

```
[ 31.,  15.,  nan,  27.],
[ nan, -41.,  nan,  30.],
[ 19.,  29.,  nan,  14.],
[ 32.,  nan,  38.,  -1.],
[-21., -31., -31., -36.],
[-11., -18.,  15., -41.],
[  7., -18., -19.,  24.],
[-27., -15.,  25.,   5.],
[-22.,  nan,  nan,  nan],
[-14.,   3., -45., -12.],
[-33.,  29., -46.,  -8.],
[  8., -19., -49.,  15.],
[ -9.,   7., -15., -39.],
[ -4.,  nan,  41.,  nan],
[-36.,  49.,  nan, -38.],
[ -8.,  34.,  25.,  18.],
[-44.,  18.,  nan, -47.],
[ 26.,   2.,  28., -35.],
[ nan,  nan,   8.,  nan],
[ 29., -37.,  35.,  -2.]]
```

2.5) Convert a1 to a dataframe named df2 and name the columns a, b, c, and d.

In [23]: #2.5

df2=pd.DataFrame(data=a1,columns=["a","b","c","d"])

df2

```
12 -27.0 -15.0  25.0   5.0
13 -22.0   NaN   NaN   NaN
14 -14.0   3.0 -45.0 -12.0

15 -33.0  29.0 -46.0  -8.0
16   8.0 -19.0 -49.0  15.0
17  -9.0   7.0 -15.0 -39.0
18  -4.0   NaN  41.0   NaN
19 -36.0  49.0   NaN -38.0
20  -8.0  34.0  25.0  18.0
21 -44.0  18.0   NaN -47.0
22  26.0   2.0  28.0 -35.0
23   NaN   NaN   8.0   NaN
```

2.6) Use mean() to return the mean value of each row. Use the axis="columns" parameter to calculate the mean value across columns.

```
In [24]: #2.6
df2.mean(axis="columns")
```

```
Out[24]: 0      2.666667
1      3.000000
2     31.666667
3     20.333333
4     23.666667
5     15.000000
6     -5.500000
7     20.666667
8     23.000000
9    -29.750000
10    -13.750000
11     -1.500000
12     -3.000000
13    -22.000000
14    -17.000000
15    -14.500000
16    -11.250000
17    -14.000000
18     18.500000
19      0.000000
```

2.7) Create a new column in df2 named group that takes on a value of 1 if the row mean is positive, and it takes on a value of 2 if the row mean is not positive.

```
In [25]: #2.7
df2['group'] = np.where(df2.mean(axis="columns") > 0, "1", "2")
df2
```

```
Out[25]:
```

| | a | b | c | d | group |
|----|-------|-------|-------|-------|-------|
| 0 | -6.0 | -3.0 | NaN | 17.0 | 1 |
| 1 | 17.0 | -41.0 | 33.0 | NaN | 1 |
| 2 | NaN | 37.0 | 20.0 | 38.0 | 1 |
| 3 | 38.0 | NaN | 8.0 | 15.0 | 1 |
| 4 | NaN | 37.0 | -4.0 | 38.0 | 1 |
| 5 | 31.0 | -13.0 | NaN | 27.0 | 1 |
| 6 | NaN | -41.0 | NaN | 30.0 | 2 |
| 7 | 19.0 | 29.0 | NaN | 14.0 | 1 |
| 8 | 32.0 | NaN | 38.0 | -1.0 | 1 |
| 9 | -21.0 | -31.0 | -31.0 | -36.0 | 2 |
| 10 | -11.0 | -18.0 | 15.0 | -41.0 | 2 |

2.8) Use groupby() to group df2 by the group column and use size() to calculate the group size

```
In [26]: #2.8
df2.groupby("group").size()
```

```
Out[26]: group
1      13
2      12
dtype: int64
```

2.9) Use notnull() and sum() to calculate the number of non-missing values in each column in df2.

```
In [27]: #2.9
df2.notnull().sum()
```

```
Out[27]: a      21
b      20
c      18
d      21
group   25
dtype: int64
```

2.10) Filter and return rows in df2 that have missing values in column c.

```
In [28]: #2.10
df2[df2["c"].isnull()]
```

```
Out[28]:
```

| | a | b | c | d | group |
|----|-------|-------|-----|-------|-------|
| 0 | -6.0 | -3.0 | NaN | 17.0 | 1 |
| 5 | 31.0 | -13.0 | NaN | 27.0 | 1 |
| 6 | NaN | -41.0 | NaN | 30.0 | 2 |
| 7 | 19.0 | 29.0 | NaN | 14.0 | 1 |
| 13 | -22.0 | NaN | NaN | NaN | 2 |
| 19 | -36.0 | 49.0 | NaN | -38.0 | 2 |
| 21 | -44.0 | 18.0 | NaN | -47.0 | 2 |

2.11) Use dropna() and the subset parameter to remove rows in df2 that have missing values in column c. Keep inplace=False by default to not modify df2 in place.

In [29]: `#2.11`
`df2.dropna(subset=["c"])`

| | | | | | |
|----|-------|-------|-------|-------|---|
| 9 | -21.0 | -31.0 | -31.0 | -36.0 | 2 |
| 10 | -11.0 | -18.0 | 15.0 | -41.0 | 2 |
| 11 | 7.0 | -18.0 | -19.0 | 24.0 | 2 |
| 12 | -27.0 | -15.0 | 25.0 | 5.0 | 2 |
| 14 | -14.0 | 3.0 | -45.0 | -12.0 | 2 |
| 15 | -33.0 | 29.0 | -46.0 | -8.0 | 2 |
| 16 | 8.0 | -19.0 | -49.0 | 15.0 | 2 |
| 17 | -9.0 | 7.0 | -15.0 | -39.0 | 2 |
| 18 | -4.0 | NaN | 41.0 | NaN | 1 |
| 20 | -8.0 | 34.0 | 25.0 | 18.0 | 1 |
| 22 | 26.0 | 2.0 | 28.0 | -35.0 | 1 |
| 23 | NaN | NaN | 8.0 | NaN | 1 |
| 24 | 29.0 | -37.0 | 35.0 | -2.0 | 1 |

2.12) Use dropna() and the thresh parameter to keep rows in df2 that have at least 4 non-missing values. Keep inplace=False by default to not modify df2 in place.

In [30]: `#2.12`
`df2.dropna(thresh=4, inplace=False)`

| | | | | | |
|----|-------|-------|-------|-------|---|
| 9 | -21.0 | -31.0 | -31.0 | -36.0 | 2 |
| 10 | -11.0 | -18.0 | 15.0 | -41.0 | 2 |
| 11 | 7.0 | -18.0 | -19.0 | 24.0 | 2 |
| 12 | -27.0 | -15.0 | 25.0 | 5.0 | 2 |
| 14 | -14.0 | 3.0 | -45.0 | -12.0 | 2 |
| 15 | -33.0 | 29.0 | -46.0 | -8.0 | 2 |
| 16 | 8.0 | -19.0 | -49.0 | 15.0 | 2 |
| 17 | -9.0 | 7.0 | -15.0 | -39.0 | 2 |
| 19 | -36.0 | 49.0 | NaN | -38.0 | 2 |
| 20 | -8.0 | 34.0 | 25.0 | 18.0 | 1 |
| 21 | -44.0 | 18.0 | NaN | -47.0 | 2 |
| 22 | 26.0 | 2.0 | 28.0 | -35.0 | 1 |

2.13) Calculate the median value of each column in df2. Use fillna() to fill missing values in df2 with the column median. Keep inplace=False by default to not modify df2 in place.

```
In [31]: #2.13
df2.fillna(df2.median(),inplace=False)
```

Out[31]:

| | a | b | c | d | group |
|----|-------|-------|-------|-------|-------|
| 0 | -6.0 | -3.0 | 11.5 | 17.0 | 1 |
| 1 | 17.0 | -41.0 | 33.0 | 5.0 | 1 |
| 2 | -6.0 | 37.0 | 20.0 | 38.0 | 1 |
| 3 | 38.0 | -0.5 | 8.0 | 15.0 | 1 |
| 4 | -6.0 | 37.0 | -4.0 | 38.0 | 1 |
| 5 | 31.0 | -13.0 | 11.5 | 27.0 | 1 |
| 6 | -6.0 | -41.0 | 11.5 | 30.0 | 2 |
| 7 | 19.0 | 29.0 | 11.5 | 14.0 | 1 |
| 8 | 32.0 | -0.5 | 38.0 | -1.0 | 1 |
| 9 | -21.0 | -31.0 | -31.0 | -36.0 | 2 |
| 10 | -11.0 | -18.0 | 15.0 | -41.0 | 2 |

2.14) Calculate the median value of each column in df2. Use fillna() to fill missing values in df2 with the column median. Keep inplace=False by default to not modify df2 in place.

```
In [32]: #2.14
df2.groupby("group").median()
```

Out[32]:

| | a | b | c | d |
|-------|-------|-------|-------|-------|
| group | | | | |
| 1 | 22.5 | 2.0 | 26.5 | 16.0 |
| 2 | -21.0 | -15.0 | -25.0 | -12.0 |

2.15) This question aims to fill missing values in df2 with the group median. Use groupby() and apply() to group df2 by the group column and apply a lambda function to all the groups. Use fillna() and median() in the lambda function to replace missing values with the group median. Take column a in df2 as an example, we want to replace missing values in group 1 with 22.5, and missing values in group 2 with -21. Keep inplace=False by default to not modify df2 in place

In [33]: **#2.15**
`df2.groupby("group").apply(lambda x:x.fillna(x.median(),inplace=False))`

Out[33]:

| | a | b | c | d | group |
|----|-------|-------|-------|-------|-------|
| 0 | -6.0 | -3.0 | 26.5 | 17.0 | 1 |
| 1 | 17.0 | -41.0 | 33.0 | 16.0 | 1 |
| 2 | 22.5 | 37.0 | 20.0 | 38.0 | 1 |
| 3 | 38.0 | 2.0 | 8.0 | 15.0 | 1 |
| 4 | 22.5 | 37.0 | -4.0 | 38.0 | 1 |
| 5 | 31.0 | -13.0 | 26.5 | 27.0 | 1 |
| 6 | -21.0 | -41.0 | -25.0 | 30.0 | 2 |
| 7 | 19.0 | 29.0 | 26.5 | 14.0 | 1 |
| 8 | 32.0 | 2.0 | 38.0 | -1.0 | 1 |
| 9 | -21.0 | -31.0 | -31.0 | -36.0 | 2 |
| 10 | -11.0 | -18.0 | 15.0 | -41.0 | 2 |
| 11 | 7.0 | -18.0 | -19.0 | 24.0 | 2 |
| 12 | -27.0 | -15.0 | 25.0 | 5.0 | 2 |
| 13 | -22.0 | -15.0 | -25.0 | -12.0 | 2 |
| 14 | -14.0 | 3.0 | -45.0 | -12.0 | 2 |
| 15 | -33.0 | 29.0 | -46.0 | -8.0 | 2 |
| 16 | 8.0 | -19.0 | -49.0 | 15.0 | 2 |
| 17 | -9.0 | 7.0 | -15.0 | -39.0 | 2 |
| 18 | -4.0 | 2.0 | 41.0 | 16.0 | 1 |
| 19 | -36.0 | 49.0 | -25.0 | -38.0 | 2 |
| 20 | -8.0 | 34.0 | 25.0 | 18.0 | 1 |
| 21 | -44.0 | 18.0 | -25.0 | -47.0 | 2 |
| 22 | 26.0 | 2.0 | 28.0 | -35.0 | 1 |
| 23 | 22.5 | 2.0 | 8.0 | 16.0 | 1 |
| 24 | 29.0 | -37.0 | 35.0 | -2.0 | 1 |